



NYU

Center for
Data Science

Week 02.1: Relations

DS-GA 1004: Big Data

This week




- **Relational databases**
- SQL
- Transactions

Database research is a huge topic!

This intro will be brief and cursory.

Database management systems (DBMS)

- DBMS's job is to provide
 - Data integrity / consistency
 - Concurrent access
 - Efficient storage and access
 - Standardized format / administration
 - Standardized query interface (language)
- DBMS come in many flavors
 - **Relational (RDBMS)**
 - Semi-structured (e.g., XML)
 - Object-oriented
 - Object-relational
 - ...



The relational model

- **High-level:** tables of data that you're probably used to
 - Spreadsheets, dataframes, numerical arrays, etc.
- Each column represents a **set** of possible values (numbers, strings, etc)

The relational model

- **High-level:** tables of data that you're probably used to
 - Spreadsheets, dataframes, numerical arrays, etc.
- Each column represents a **set** of possible values (numbers, strings, etc)
- A **relation** over sets $A_1, A_2, A_3, \dots, A_n$ is a **subset** of their cartesian product
 - $R \subseteq A_1 \times A_2 \times A_3 \times \dots \times A_n$
 - The **rows** of the table are elements of R , also known as **tuples**
 - $(a_1, a_2, \dots, a_n) \in R \Rightarrow a_1 \in A_1, a_2 \in A_2, \dots, a_n \in A_n$

Example: dinosaurs

- $A_1 = \{s \mid s \text{ is a string}\}$
 $A_2 = \{\text{"Jurassic", "Cretaceous", "Devonian", "Triassic", ...}\}$
 $A_3 = \{\text{"Carnivore", "Herbivore", "Omnivore", ...}\}$
 $A_4 = \{\text{False, True}\}$
- Any A_i could be finite or infinite
- $R \subseteq A_1 \times A_2 \times A_3 \times A_4$ need not contain all combinations!

Species	Era	Diet	Awesome
T. Rex	Cretaceous	Carnivore	True
Stegosaurus	Jurassic	Herbivore	True
Ankylosaurus	Cretaceous	Herbivore	False

Aside: why “relations” and not “tables”?

- Relations are the abstract model of data
- **Table** refers to an **explicitly** constructed relation
 - I.E., records you’ve observed / collected
- Other relations in a DB:
 - **view**: a relation defined **implicitly**, and constructed **dynamically** at run-time
 - **temporary table**: the output of a **query**

Properties of relations

- $R \subseteq A_1 \times A_2 \times A_3 \times \dots \times A_n$ is a set
 - The tuples (rows) of R are **unordered**
 - Tuples are **unique** \Rightarrow no duplicates!
 - Relations over common domains (columns) can be combined by set operations

Species	Era	Diet	Awesome
T. Rex	Cretaceous	Carnivore	True
Stegosaurus	Jurassic	Herbivore	True
Ankylosaurus	Cretaceous	Herbivore	False

Properties of relations

- $R \subseteq A_1 \times A_2 \times A_3 \times \dots \times A_n$ is a set
 - The tuples (rows) of R are **unordered**
 - Tuples are **unique** \Rightarrow no duplicates!
 - Relations over common domains (columns) can be combined by set operations
- In practice, add a column (e.g., A_0) with **identifiers** to force uniqueness
 - This is not (usually) part of the data, but is generated automatically by the DBMS
 - ID fields are often used as **primary keys**, and give a default order to rows

id	Species	Era	Diet	Awesome
1	T. Rex	Cretaceous	Carnivore	True
2	Stegosaurus	Jurassic	Herbivore	True
3	Ankylosaurus	Cretaceous	Herbivore	False

Schemas

- A relation is defined by a **schema**:

id	Species	Era	Diet	Awesome
1	T. Rex	Cretaceous	Carnivore	True
2	Stegosaurus	Jurassic	Herbivore	True
3	Ankylosaurus	Cretaceous	Herbivore	False
4	Homer	Boomer	Donuts	False

`Dinosaur(id: int, Species: string, Era: string, Diet: string, Awesome: boolean)`

- Any tuple `(int, string, string, string, boolean)` is valid under this schema
 - \Rightarrow Schemas enforce type (syntax), but not semantics!



Relational databases

- A relational database consists of one or more relational schemas
- Structured data can be encoded by **joining** on **shared attributes**
- The collection of schemas defines your **data model**

id	Species	Era	Diet	Awesome
1	T. Rex	Cretaceous	Carnivore	True
2	Stegosaurus	Jurassic	Herbivore	True
3	Ankylosaurus	Cretaceous	Herbivore	False



id	Name	Species	Internals
1	Earl Sinclair	Megalosaurus	Puppet
2	Grimlock	T. Rex	Robot
3	Snarl	Stegosaurus	Robot

Keys


- Keys are what determine the **identity** of a row
- Keys can be simple (single column) or **compound** (two or more columns)
 - Example: (First Name, Last Name)
 - This prevents two rows with the same combination of first and last name
- You can have **primary** and **alternate keys**
 - Usually a good idea to keep a primary numeric key as well as others you may want...

id	First Name	Last Name	Age
1	Homer	Simpson	39
2	Marge	Simpson	39
3	Bart	Simpson	10
4	Homer	Thompson	39
5	Homer	Simpson	28

Foreign Keys

- A **key** from one relation can be a **column** in another
 - This is called a **FOREIGN KEY** constraint
- This can be used to ensure reference consistency **between** tables/relations
- **This is not automatic**: must be included in the **schema** definition!

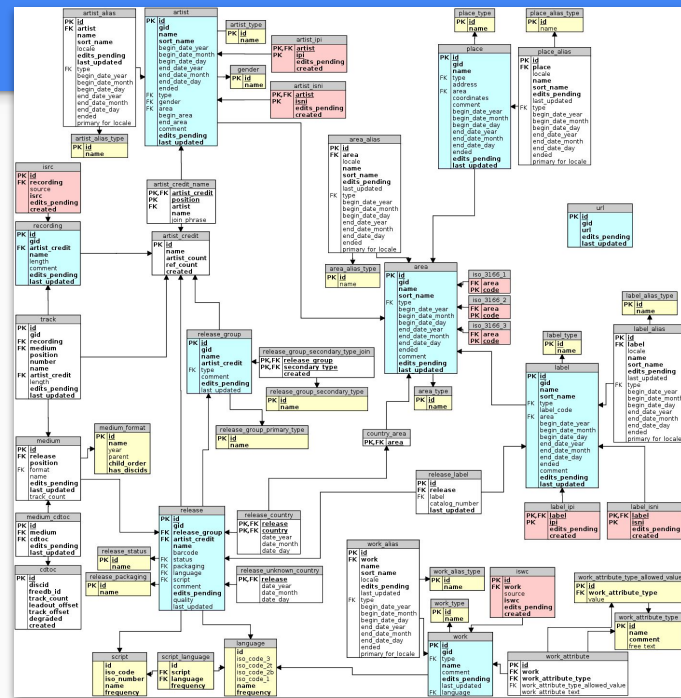
id	Species	Era	Diet	Awesome
1	T. Rex	Cretaceous	Carnivore	True
2	Stegosaurus	Jurassic	Herbivore	True
3	Ankylosaurus	Cretaceous	Herbivore	False



id	Name	DinosaurID	Internals
1	Earl Sinclair	25	Puppet
2	Grimlock	1	Robot
3	Snarl	3	Robot

Normalization

- A database schema is **normal** if data is not redundantly stored
 - Use **identifiers**, not **values**, to link between relations
- Modifying a record is easy if it exists in exactly one place
- But it can also be difficult
 - Reading complex data can be cumbersome
 - Multiple levels of indirection



Summary

Relations are cool!

- We use relational data every day without thinking of it
- Databases consist of one or more relations
- Putting data into a relational model can make it easier to work with
- Schemas provide some degree of safety and validation